

Parallel Algorithms and Programming

Lab 5 report

Member:
Son-Tung DO
Johana MARKU

1 Fox Algorithm

We successfully implemented the Fox's algorithm with MPI.

Since we do not handle the data distribution in this exercises, all the processes are initialized with a value a and b as a block of matrix A and B. Also a variable c to store the result of computation of the final matrix in the same block.

Given the number or process available N , we always create a matrix of size $n \times n$ where $n = \sqrt{N}$ if $N \bmod 2 = 0$ or $\sqrt{N-1}$ if $N \bmod 2 \neq 0$

We create a grid process, and 2 sub-communicator are for row and column. The row communicator is used for broadcasting the value of matrix A as every stage, the whole row using the same value of a certain block in that row. The column matrix is used for shifting the value of matrix B.

N.B: The code provided has the matrix value initialized in each block corresponded to the same rank of the process holding it for the purpose of testing, the code for random value also provided

2 Cannon Algorithm

We successfully implemented the algorithm with MPI.

The same setup is used from the previous exercise: Matrix value initialization, Grid size, sub-communicator. The different comes from the algorithm itself, when we have to do a pre-skewing and post-skewing stage before and after the loop run.

There are optimization we did in this exercise, we created more local variables to store values of matrix A and B during the shifting. Because of that, we avoid modifying the original value of matrix A and B. Thanks to this, the post-skewing stage is unnecessary. This will save us some communication cost to restore the initial state of the matrix.

N.B: The optimization is done in file ex2-op.c

3 Data distribution

For this exercises, we tried to utilized the vector datatype of MPI by define a typeMatCol, and type MatRow. While doing the Scatterv function, we stuck dealing with distributing the MatB while MatA is run perfectly.

The only thing, we noticed is that the blocks of value that got problem are all supposed to have the same data. We have that problem and don't know how to fix it. We did try to used both dynamic allocation and static allocation.

4 Performance Evaluation

We did some performance test with Fox Algorithm, Cannon Algorithm, also with a sequential code (written by ourselves)

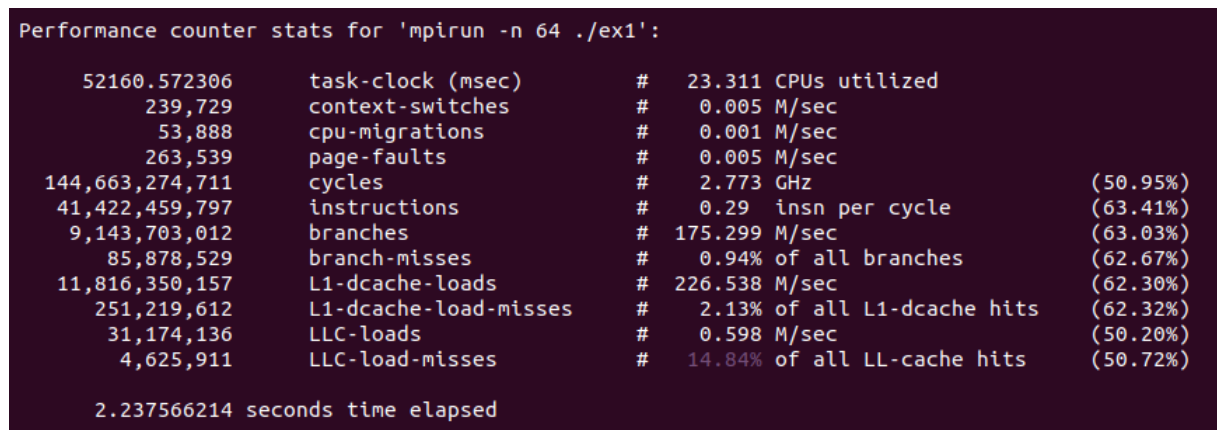


Figure 1: Fox algorithm with 8x8 matrix

```

Performance counter stats for 'mpirun -n 64 ./ex2':

 58236.937831    task-clock (msec)    #    24.493 CPUs utilized
   292,254      context-switches        #    0.005 M/sec
    54,002      cpu-migrations          #    0.927 K/sec
   263,832      page-faults             #    0.005 M/sec
161,953,265,688  cycles                    #    2.781 GHz                    (50.62%)
 46,264,153,570  instructions            #    0.29  insn per cycle          (63.13%)
 10,329,897,081  branches                 #   177.377 M/sec                  (62.85%)
   90,991,829    branch-misses              #    0.88% of all branches         (62.73%)
 13,356,160,516  L1-dcache-loads          #   229.342 M/sec                  (62.54%)
 280,486,560     L1-dcache-load-misses       #    2.10% of all L1-dcache hits    (62.42%)
   32,486,627    LLC-loads                #    0.558 M/sec                   (50.18%)
    5,067,407    LLC-load-misses           #   15.60% of all LL-cache hits     (50.46%)

 2.377730854 seconds time elapsed

```

Figure 2: Cannon algorithm with 8x8 matrix

```

Performance counter stats for 'mpirun -n 64 ./ex2-op':

 56279.062778    task-clock (msec)    #    23.791 CPUs utilized
   318,358      context-switches        #    0.006 M/sec
    55,810      cpu-migrations          #    0.992 K/sec
   263,474      page-faults             #    0.005 M/sec
156,844,996,423  cycles                    #    2.787 GHz                    (50.27%)
 42,703,411,774  instructions            #    0.27  insn per cycle          (62.96%)
   9,430,690,073  branches                 #   167.570 M/sec                  (62.93%)
   92,212,928    branch-misses              #    0.98% of all branches         (62.90%)
 12,310,408,203  L1-dcache-loads          #   218.739 M/sec                  (62.80%)
 304,046,317     L1-dcache-load-misses       #    2.47% of all L1-dcache hits    (62.81%)
   29,167,822    LLC-loads                #    0.518 M/sec                   (49.96%)
    4,780,283    LLC-load-misses           #   16.39% of all LL-cache hits     (50.28%)

 2.365564677 seconds time elapsed

```

Figure 3: Cannon algorithm with 8x8 matrix (optimized version)

```

Performance counter stats for 'mpirun -n 64 ./seq':

 175.417895      task-clock (msec)    #    0.362 CPUs utilized
     409         context-switches        #    0.002 M/sec
     134         cpu-migrations          #    0.764 K/sec
    19,608       page-faults             #    0.112 M/sec
 533,461,438     cycles                    #    3.041 GHz                    (52.08%)
 604,911,838     instructions            #    1.13  insn per cycle          (68.31%)
 135,044,141     branches                 #   769.842 M/sec                  (69.04%)
   2,772,900     branch-misses              #    2.05% of all branches         (79.44%)
 171,242,657     L1-dcache-loads          #   976.198 M/sec                  (76.67%)
 11,774,676      L1-dcache-load-misses     #    6.88% of all L1-dcache hits    (71.33%)
   1,414,084     LLC-loads                #    8.061 M/sec                   (54.93%)
    238,117      LLC-load-misses           #   16.84% of all LL-cache hits     (48.05%)

 0.484986967 seconds time elapsed

```

Figure 4: Sequential with 8x8 matrix

We realized that our implementation has a very simple local computation as 1 value is hold in one process. This make the sequential time so small as they don't have communication cost.

There is one remark is our Canon optimized version actually run faster than the normal implementation. But since the size of the matrix that we test is not very big, the amount is not much.